

# COL331 Assignment 2 Report

Submitted by:

Dharmeswar Basumatary (2020CS50423)

Mohd Bilal Bin Rehmat (2020CS50431)

Attributes added in struct proc in the proc.h file are:

```
int exec_time;

int arrival_time;

int deadline;

int elapsed_time;

int sched_policy;

//For RMS

int rate;

int weight;
```

## System calls:

We are required to add the signature of system calls in user.h,defs.h

In usys.S SYSCALL(sched\_policy), SYSCALL(exec\_time), SYSCALL(deadline), SYSCALL(rate)

1) Sys\_sched\_policy();

It is implemented in sysproc.c file.

It parses the arguments pid and policy using argint(int, int\*) function and calls the \_sched\_policy(int pid,int policy) defined in proc.c

In \_sched\_policy() function:

Scaning the ptable to find the structure of the process p using pid.

If the policy is 0: (EDF)

Checking for the schedulability test  $U \leq 1$ , where  $U = \text{Sum of (exec\_time/deadline) of all the process to be scheduled.}$

To implement this, a global variable U is used which store sum of exec\_time/deadline for all the process with sched\_policy =0.

If  $U(\text{global variable}) + p \rightarrow \text{exec\_time} / p \rightarrow \text{deadline} \leq 1$  then  $p \rightarrow \text{sched\_policy}$  is set to 0 else the process is killed.

If the policy is 1:(RMS)

Checking for schedulability test for RMS:

$$U \leq n * (2^{1/n} - 1)$$

N is the no of process. A global variable  $n\_rms\_procs$  stores the no of process with  $\text{sched\_policy} = 1$ .

$$U = \text{SUM} (p \rightarrow \text{exec\_time} * p \rightarrow \text{rate} / 100) \text{ (here)}$$

Because rate is ticks/second and 1 tick = 10 milliseconds.

Helper function  $\text{powr}(\text{base}, n)$ :

It returns  $\text{base}^{1/n} - 1$ ;

Global variable  $\text{RMS\_U}$  is used which stores Utility of all the processes with  $\text{sched\_policy} = 1$ .

If  $\text{RMS\_schedulability}$  for the process fails, then the process is killed.

Code Snippet:

```
12 // Setting sched_policy
13 int _sched_policy(int pid, int policy)
14 {
15     struct proc *p;
16     acquire(&ptable.lock);
17     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
18     {
19         if (p->pid == pid)
20         {
21             break;
22         }
23     }
24     if (policy == 0)
25     {
26         double temp = U + ((double)p->exec_time / (double)p->deadline);
27         if (temp > 1.0)
28         {
29             release(&ptable.lock);
30             kill(pid);
31             return -22;
32         }
33         p->sched_policy = policy;
34         p->arrival_time = ticks;
35         U = temp;
36         release(&ptable.lock);
37         return 0;
38     }
39     else
40     {
41         // rms_policy = policy;
42         double t1 = RMS_U + ((double)(p->exec_time * p->rate) / 100.0);
43         double bound = powr((double)2, n_rms_proc+1);
44         bound = bound*(double)(n_rms_proc+1);
45         if (t1 > bound){
46             release(&ptable.lock);
47             kill(pid);
48             return -22;
49         }
50         // cprintf("%s %d\n", "setpolicy: ", n_rms_proc);
51         p->sched_policy = policy;
52         p->arrival_time = ticks;
53         n_rms_proc++;
54         RMS_U = t1;
55         release(&ptable.lock);
56         return 0;
57     }
58 }
59 }
```

## 2) Sys\_exec\_time();

It is implemented in sysproc.c file.

It parses the arguments pid and exec\_time using argint(int, int\*) function and calls the \_exec\_time(int pid,int exec\_time) defined in proc.c

In \_exec\_time():

Scanning the ptable to find the structure of the process p using pid.

Then, p->exec\_time = exec\_time is set(if not set earlier).

Code Snippet:

```
0 // Setting exec_time
1 int _exec_time(int pid, int exec_time)
2 {
3     struct proc *p;
4     acquire(&ptable.lock);
5     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
6     {
7         if (p->pid == pid)
8         {
9             if (p->exec_time != -1)
10            {
11                release(&ptable.lock);
12                return -22;
13            }
14            break;
15        }
16    }
17    p->exec_time = exec_time;
18    release(&ptable.lock);
19    return 0;
20 }
21
```

## 3) Sys\_deadline();

It is implemented in sysproc.c file.

It parses the arguments pid and deadline using argint(int, int\*) function and calls the \_deadline(int pid,int deadline) defined in proc.c

In \_deadline():

Scanning the ptable to find the structure of the process p using pid.

Then, p->deadline = deadline is set (if not set earlier).

Code Snippet:

```

171
172 // Setting deadline
173 int _deadline(int pid, int deadline)
174 {
175     struct proc *p;
176     acquire(&ptable.lock);
177     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
178     {
179         if (p->pid == pid)
180         {
181             if (p->deadline != -1)
182             {
183                 release(&ptable.lock);
184                 return -22;
185             }
186             break;
187         }
188     }
189     p->deadline = deadline;
190     release(&ptable.lock);
191     return 0;
192 }
193

```

#### 4) Sys\_rate();

It is implemented in sysproc.c file.

It parses the arguments pid and rate using argint(int, int\*) function

and calls the \_rate(int pid,int rate) defined in proc.c

Scanning the ptable to find the structure of the process p using pid.

P->rate = rate is set. (if not set earlier).

This function also sets the weight of the process.

P->weight = Ceiling((30-rate) \*3/29)).

Code Snippet:

```

// Setting rate for rms Scheduler
int _rate(int pid, int rate)
{
    struct proc *p;
    acquire(&ptable.lock);
    for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
    {
        if (p->pid == pid)
        {
            if (p->rate != -1)
            {
                release(&ptable.lock);
                return -22;
            }
            break;
        }
    }
    p->rate = rate;
    int wt = ((30 - rate) * 3);
    if (wt < 29)
    {
        p->weight = 1;
    }
    else
    {
        if (wt % 29 == 0) p->weight = wt/29;
        else p->weight = wt/29 + 1;
    }
    release(&ptable.lock);
    return 0;
}

```

## Implementation of Scheduling Algorithms:

Both the scheduling algorithms **EDF (Earliest Deadline First)** and **RMS (Rate Monotonic Scheduling)** are implemented in the **scheduler()** function in the **proc.c** along with the original xv6 round robin algorithm.

Schedulability test both EDF And RMS is done during the **sys\_sched\_policy()** system call.

If incoming process is not schedulable it is killed in the system call.

### EDF (Earliest Deadline First):

#### Implementation Idea:

Find a EDF schedulable process with minimum deadline+ arrival time from the ptable. If two or more processes have the same value of deadline + arrival time , pick the one with minimum pid and schedule it.

#### Implementation:

Helper functions:

`int get_policy();`

It scans the process table and returns an integer 0 or 1 or -1;

If there is a EDF schedulable process it returns 0. If there is a RMS schedulable process it returns 1. If there is both EDF and RMS schedulable processes in the ptable it returns either 0 or 1.(no priority set between EDF and RMS).

If there is no RMS or EDF schedulable process it returns -1;

Code snippet:

```
9
10 int get_policy()
11 {
12     for (int i = 0; i < NPROC; i++)
13     {
14         struct proc temp = ptable.proc[i];
15         if (temp.state != RUNNABLE)
16             continue;
17         if (temp.sched_policy != -1)
18         {
19             return temp.sched_policy;
20         }
21     }
22     return -1;
23 }
```

`int Find_edf_ind();`

It scans the process table and returns the index of a runnable process with minimum deadline+arrival time and minimum pid(if more than one process has same deadline + arrival time) in the process table.

## Code snippet:

```
for (int i = 0; i < NPROC; i++)
{
    struct proc temp = ptable.proc[i];
    if (temp.state == RUNNABLE && temp.pid != 0 && temp.sched_policy == 0)
    {
        if (start == 0)
        {
            start = 1;
            index = i;
            min = temp.deadline + temp.arrival_time;
            continue;
        }
        if (temp.deadline + temp.arrival_time < min || (temp.deadline + temp.arrival_time == min && ptable.proc[index].pid > temp.pid))
        {
            min = temp.deadline + temp.arrival_time;
            index = i;
        }
    }
}
return index;
}
```

In the Scheduler () function in proc.c

Schedule\_policy = Get\_policy() function is called.

If its return value is 0 then, then find\_edf\_ind() function is called to get the index of the process in the ptable to be scheduled using EDF. Then context switch is done. At the time of context switch elapsed\_time of the process is incremented by one (in ticks);

Code Snippet for edf part in scheduler:

```
else if (schedule_policy == 0)
{
    int ind = find_edf_ind();
    p = &ptable.proc[ind];
    p->elapsed_time++;
    c->proc = p;
    switchvm(p);
    p->state = RUNNING;

    swtch(&(c->scheduler), p->context);
    switchkvm();

    // Process is done running for now.
    // It should have changed its p->state before coming back.
    c->proc = 0;
}
```

## RMS (Rate Monotonic Scheduling)

### Implementation Idea:

Find a RMS schedulable process from the ptable with minimum weight (where weight =

$\text{Ceiling}((30 - \text{rate}) * 3 / 29)$ ). If two or more processes have same weight pick the process with smallest pid).

Implementaion:

Weight of each process is stored in weight attribute of the process when sys\_rate (pid,rate) system call is called.

Helper functions:

Int find\_rms\_ind();

It scans the process table and returns the index of a runnable process with minimum weight and minimum pid(if more than one process has same minimum weight) from the process table.

Code Snippet:

```
for (int i = 0; i < NPROC; i++)
{
    struct proc temp = ptable.proc[i];
    if (temp.state == RUNNABLE && temp.pid != 0 && temp.sched_policy == 1)
    {
        if (start == 0)
        {
            start = 1;
            index = i;
            weight = temp.weight;
            continue;
        }
        if (temp.weight < weight || (temp.weight == weight && ptable.proc[index].pid > temp.pid))
        {
            weight = temp.weight;
            index = i;
        }
    }
}
return index;
```

In the Scheduler () function in proc.c

Schedule\_policy = Get\_policy() function is called.

If the return value of get\_policy is 1 then, then find\_rms\_ind() function is called to get the index of the process in the ptable to be scheduled using RMS. Then context switch is done. At the time of context switch elapsed\_time of the process is incremented by one (in ticks);

Code Snippet for RMS in scheduler function:

```
}  
else if (schedule_policy == 1)  
{  
    int rm_ind = find_rms_ind();  
    p = &ptable.proc[rm_ind];  
    p->elapsed_time++;  
    c->proc = p;  
    switchvm(p);  
    p->state = RUNNING;  
  
    swtch(&(c->scheduler), p->context);  
    switchkvm();  
    c->proc = 0;  
}  
release(&ptable.lock);  
}
```